

When the Work Is Gone

On Meaning After the Income Question Is Solved

Cloud Wu

May 2026

Contents

Prologue — The Wrong Question	1
Chapter 1 — What Work Was Secretly Doing	4
Chapter 2 — Money Can't Fill the Hole	8
Chapter 3 — Are You Just Catastrophising?	12
Chapter 4 — The First to Hit the Wall	15
Chapter 5 — How Badly Can This Go?	20
Chapter 6 — Finding Your Place Again	24
Chapter 7 — What Policy Can Do	29
Epilogue — Living Honestly With Uncertainty	33
References	36

Prologue — The Wrong Question

For about a decade now, a single question has organised almost everything written about artificial intelligence and the future: *Will it take my job?*

It is a reasonable question. It is also, I think, the wrong one — or at least, the one that keeps us from seeing the harder thing coming up behind it.

Let me start with a place where the future has, in a small way, already arrived.

Stack Overflow is the website where the world's programmers go to ask and answer technical questions. For fifteen years it was one of the load-bearing pillars of how software actually gets made: you hit a problem, you searched, and someone — often a stranger, for free, years earlier — had written the answer. It ran on a strange and beautiful economy. Nobody was paid. People answered questions for reputation, for the pleasure of solving a hard problem, for membership in a community of people who cared about the same obscure things.

Then, at the end of November 2022, ChatGPT arrived, and within a year it could answer most of those questions instantly, privately, without making you feel stupid for asking.

What happened next is not a forecast. It is a measurement. By my own pull of Stack Overflow's public data, the number of questions and answers posted each month fell roughly 45% in the first year — and by early 2026, around 95%. Not a decline. A collapse. An entire ecosystem of voluntary human contribution, drained in three years.

I spent weeks studying that collapse in detail — who stopped, who stayed, and why. (That study, and the ways it proved me wrong, is a thread that runs through this book; I'll come back to it.) But the single most important thing I learned from it had nothing to do with programmers.

It was this. When the work went away, the *income* of those people was mostly fine — they were never paid for it anyway. What vanished was something else. The reason to show up. The place where a certain kind of person had been competent, recognised, needed, connected. The platform didn't take their salaries. It took the thing the work had been *giving* them that they hadn't known to name.

And that is the question this book is actually about. Not *will AI take your job* — but:

When the work is gone, even if the money is somehow fine, what is left?

Two debates, and the one nobody is having

Strip away the noise and the public conversation about AI and work is really two debates.

The loud one is **economic**. Will there be enough jobs? Can we afford universal basic income, or can we not afford to skip it? Whose wages rise, whose collapse? This debate has economists, think tanks, central banks, and an ocean of data. It matters. It is also not my subject, and I am not the person to add to it.

The quiet one — the one with almost no one in the room — is about **meaning**. Suppose the economic debate resolves in the most optimistic way imaginable: new jobs appear, or a generous basic income arrives, and nobody goes hungry. *Even then*, a question remains that money does not touch. Work has never been only a paycheck. It has been, for most adults in the modern world, the default supplier of structure, status, competence, belonging, and the sense of being needed. If AI dissolves work as the place those things come from — even while keeping everyone fed — does the society on the other side stay psychologically and socially whole?

This book lives entirely in that second, quieter debate. I am going to mostly *grant* the optimists their economics, precisely so we can look clearly at the thing their optimism does not cover.

A promise about honesty

I want to make one promise up front, because it is the only thing that makes a book like this worth your time in a genre this crowded with confident nonsense.

The honest truth is that no one knows how this goes. The future of AI is not knowable the way next year's weather is roughly knowable. But "we can't predict it" is not the same as "anything goes," and it is certainly not an excuse to either panic or pretend.

So throughout this book I will hold myself to a discipline of marking exactly what kind of claim I am making, sentence by sentence:

- **What we actually know** — backed by real data, the kind a future study could prove wrong. (There is less of this than anyone would like, and I will not pad it.)
- **What we can reasonably infer** — by analogy and mechanism, always with the places the reasoning breaks named out loud.
- **What we have to choose** — the questions that are not predictions at all, but values and designs, where the honest standard is not “is it true” but “does it hold up across many possible futures.”

When I show you data, I’ll tell you precisely what it does and does not show. When I reason from history, I’ll point at where the analogy snaps. When I’m urging you toward something, I’ll own it as my judgement, not disguise it as a finding. And when the evidence has embarrassed me — as it has, more than once — I’ll tell you that too.

This is also a book with two readers in mind, who turn out to be the same reader.

The first is the professional watching it happen to their own craft right now — the programmer, the translator, the illustrator, the copywriter, the analyst — the people standing closest to the edge as AI moves into work they spent years learning to do well. If that’s you, this book is trying to give you something more useful than reassurance or alarm: a little help, honestly come by, as you go about the work of finding your place again.

The second reader is everyone, eventually. Because the professionals are not a special case. They are the *forerunners*. What is happening to them first is, I’ll argue, a preview of a question the rest of us will each have to answer in our own time: when what you produce no longer defines you — because something else can produce it — who are you, and where does your life get its meaning?

I can’t answer that question for you; no honest book could. But I can walk a little of the way with you — clear about what we know, what we’re only guessing, and what each of us has to choose for ourselves. That’s the whole offer.

We’ll start where the loss is easiest to underestimate: with what work was quietly doing for us all along.

Chapter 1 — What Work Was Secretly Doing

In 1933, in a small Austrian village called Marienthal, a young social scientist named Marie Jahoda watched what happens to people when work disappears.

Marienthal was a company town. Almost everyone worked at the textile factory, and during the Great Depression the factory closed. Here was a chance to see something that is normally tangled up with a dozen other things: not poverty alone, not illness, not war, but the specific effect of *losing the job itself*. Jahoda and her colleagues did not just count incomes. They measured how fast people walked across the village square. They read diaries. They noted how often the library books were borrowed, how meals drifted later and later, how men who had nothing to do nonetheless could not say where their day had gone.

What they found was not simply that people were poorer. It was that people came *undone*. Time itself dissolved. A villager wrote that he no longer knew what day it was. Political and social life withered. People had more free hours than ever and used them for less. The unemployed of Marienthal were not relaxing into a hard-won leisure. They were disintegrating.

Jahoda spent the next half-century working out why — and her answer is the foundation this whole book is built on. So it's worth stating carefully.

The paycheck is the smallest part

Jahoda's insight was that a job gives you two completely different kinds of thing.

The first is obvious: money. She called this the **manifest** function of work — the reason we say we go to work.

But a job, she argued, also quietly delivers five other things — the **latent** functions — that we almost never notice precisely because employment hands them to us automatically, bundled in, whether we want them or not:

1. **Time structure.** A job carves the formless day into a shape. Morning,

tasks, lunch, more tasks, evening, weekend. Remove it and, as Marienthal showed, time doesn't become free — it becomes a swamp.

2. **Social contact.** Work throws you among people who are not your family — colleagues, clients, the regulars. A surprising share of adult human contact is simply a by-product of showing up to work.
3. **Collective purpose.** A job ties your small daily actions to something larger than yourself — a product, a mission, a customer served. You are part of an effort.
4. **Status and identity.** “What do you do?” is, in modern life, very close to “who are you?” A job is a socially legible answer to the question of your place in the world.
5. **Regular activity.** Work demands that you exercise skills, that you *do* — and being habitually competent at something turns out to be a quiet, constant source of self-respect.

Read that list again and notice something uncomfortable: **not one of those five is money, and not one of them is automatically replaced by money.** A check in the mail restores the manifest function and none of the latent ones. This is the hinge the entire book turns on, so I'll put it plainly: *if you replaced a laid-off worker's income in full, you would have addressed the smaller of the two things their job was doing for them.*

This is not just a beautiful old theory

Here is where I have to be careful, because a 1933 village study, however haunting, is in the end a single vivid story — suggestive, not yet proof. The reason Jahoda's framework counts as something we actually *know*, rather than something I'm asking you to take on faith, is what happened in the ninety years after.

Her “latent deprivation” model became one of the most tested ideas in the psychology of work. And the large meta-analyses — pooling hundreds of studies across decades and countries (Paul & Moser's 2009 analysis is the landmark) — converge on two findings that matter enormously here:

- **Unemployment damages mental health, robustly and causally.** Losing a job significantly worsens depression and anxiety; finding one significantly improves them. This is not a soft correlation; the before-and-after, lose-it-then-regain-it pattern shows up again and again.
- **The damage is mostly *not* about the money.** When researchers separate the financial blow from the loss of the latent functions, the **non-financial** losses account for the *larger* share of the harm. People are hurt by unemployment well beyond, and largely apart from, the hit

to their bank account.

That second finding is the empirical spine of this chapter, so let me be exact about what it does and doesn't say. What the evidence supports is this: across the existing literature, the psychological harm of unemployment is driven more by the loss of structure, contact, purpose, status, and activity than by the loss of income. It does *not* say money doesn't matter (it does), and it is built on *involuntary* unemployment in societies organised around work — which, I'll argue later, is both its strength and a limit when we extrapolate to an AI future that might look different.

But within those bounds, the finding is about as solid as social science gets. Work was doing five things for you that you probably never credited it with. And a society now contemplating the removal of work — even a society humane enough to keep everyone's income whole — is contemplating the removal of those five things, with no automatic plan to replace them.

Why this should unsettle the optimist and console you

For the optimist who says "AI will make us all rich and free," this chapter is the first crack in the story. Rich, maybe. But "free" assumes the thing being removed was a burden. Marienthal says the thing being removed was also a scaffold — and people fell down when it went.

But there's a more personal reading, and it's the one I want you to carry out of this chapter.

If you have felt, lately, a low and hard-to-name disquiet about your work — not "I might lose my income" but something vaguer, a sense that the ground is shifting under the part of your life you thought was solid — this chapter is permission to take that feeling seriously. It is not fragility or ingratitude. You are sensing, correctly, that your work has been quietly supplying you with structure, contact, purpose, status, and the daily exercise of competence — and that *these*, not just your salary, are what's now in question.

That sounds like bad news, and partly it is. But naming a loss precisely is the first move in doing something about it. You cannot rebuild what you cannot see. Jahoda's five functions are, in a sense, a *shopping list*: if work stops supplying them, these are the five things a life — and a society — will need to source elsewhere. The rest of this book is, in large part, about where else they might come from.

First, though, we have to clear away the most popular wrong answer —

the one that says we can simply buy our way out. That's the next chapter.

Sources for this chapter's empirical claims: Jahoda (1982), the Marienthal study (Jahoda, Lazarsfeld & Zeisel, 1933); Paul & Moser (2009), "Unemployment impairs mental health: Meta-analyses," Journal of Vocational Behavior; the latent deprivation meta-analytic literature. See ../../evidence_base.md.

Chapter 2 — Money Can't Fill the Hole

There is a comforting answer to everything in the last chapter, and almost everyone reaches for it. It goes: *fine — if AI takes the work but creates enormous wealth, we just give people the money. Universal basic income. Problem solved.*

I used to find this answer persuasive too. It is humane, it is simple, and it has serious people behind it. It is also, on the specific question of *meaning*, almost certainly wrong — and unusually for this book, we don't have to guess. This is one of the rare places where we can actually check, because the experiment has been run, more than once, on real people.

So this is the most evidence-heavy chapter in the book — the one where I'm leaning hardest on real data rather than reasoning. Let me walk through what the money actually bought.

What the cash experiments found

Finland, 2017–2018. The Finnish government took 2,000 unemployed people and simply paid them €560 a month, no strings — no job-search requirement, no clawback if they found work. Then they compared this group to everyone else. The results, published in 2020, are quietly devastating to the simple optimism:

- On *employment* — the thing the policy was loudest about — almost nothing happened. The difference in days worked was on the order of half a day in the first year.
- On *wellbeing*, though, there was a real effect. Recipients reported better mental health, less depression and loneliness, more confidence in the future, a stronger sense of economic security.

Read quickly, that second bullet sounds like victory for the money. Read carefully, it is the opposite of the thing we're chasing. What improved was **relief and security** — the easing of the grinding stress of scarcity. That is real, and it matters, and I don't want to wave it away. But it is not the same as *purpose*. Nobody in the Finnish data was reported to have found that their life now *mattered* more, that they had a reason to get up. They were less anxious. They were not, in the sense Jahoda would recognise, re-supplied with the latent functions of work.

The United States, 2020–2023. Then came the largest test yet: a study funded through Sam Altman's OpenResearch gave 1,000 low-income people \$1,000 a month for three years, against a control group, and tracked them closely. If anything could show money buying a better life, a thousand dollars a month for three years should.

The results, released in 2024, were sobering. The money went, sensibly, to essentials — rent, food, transport. And on mental health, here is the crucial finding: stress and psychological distress **did** improve — but mostly in **year one**, and the improvement had **faded by year two**.

Now, the honest part — and this is exactly the kind of moment this book exists for. It would be easy to package “the benefit faded” as proof that money's comfort is shallow and temporary. But the researchers themselves point to two confounds I'm obliged to pass on to you: year one overlapped with extra pandemic relief payments, which may have inflated the early effect, and inflation then steadily ate the purchasing power of a fixed \$1,000. So even this finding comes fenced. The *fade* might be about COVID and inflation, not about some iron law that money's comfort always evaporates. The data is real; what it proves is narrower than the tidy story you could spin from it. I'd rather hand you the fenced truth than the clean lie.

The people who got rich, and the people who stopped working. Two more bodies of evidence point the same way, from completely different directions:

- *Lottery winners.* When researchers study people who win large sums (the best work uses lottery data precisely because winning is random, which makes it a clean natural experiment), the pattern is consistent: big money produces a modest, lasting bump in *life satisfaction* — and strikingly little improvement in day-to-day mental health. And many winners keep working. Handed the means to never work again, large numbers of people choose to keep doing it anyway.
- *Retirement.* This is the closest thing we have to a preview of “money is handled, work is gone,” and it carries a sharp warning. Retirement is not uniformly good or bad for people. For some — especially those escaping grinding or degrading jobs — it lifts wellbeing. For others it brings depression and faster cognitive decline. What separates the two is not money. It is whether the person rebuilds *structure, social contact, and purpose* — Jahoda's latent functions — somewhere outside the old job. The retirees who do well are the ones who replace what work was secretly providing. The ones who don't, decline.

The finding, stated as narrowly as it deserves

Put it all together and resist, hard, the urge to overstate it:

Money reliably buys relief from material stress. There is no good evidence that it buys meaning.

That is a deliberately small claim, and its smallness is the point. I am *not* saying money doesn't matter — it plainly does; scarcity is its own kind of misery and easing it is a real good. I am *not* saying UBI is a bad idea — by this evidence it's an excellent floor. I am saying something narrower and more durable: **the hole that Chapter 1 described — the missing structure, contact, purpose, status, and competence — is not a hole that cash fills.** Across every natural experiment we have, money tops up the manifest function of work and leaves the five latent ones empty.

What this means for the optimist — and for you

For the policy optimist, this reframes UBI entirely. Cash is a **floor, not an engine**. It can prevent destitution; it cannot manufacture meaning. A post-work society that gets the money right and stops there will have solved the smaller problem and left the larger one untouched — will have funded Marienthal without curing it. (What the engine might actually be is the subject of the last third of this book.)

But here is the personal reading, and I think it's genuinely good news disguised as bad.

If some part of you has been waiting for a number — a salary, a windfall, a “financial independence” figure, a UBI — to arrive and finally make you whole, the evidence is gently telling you to stop waiting. Not because the number won't come, but because *it was never going to do that job*. The thing you actually want — to feel competent, needed, connected, part of something, located in your days — is not, it turns out, for sale. That sounds like a loss. It is actually a release: it means your meaning was never hostage to your bank balance in the first place. It is sourced elsewhere. And “elsewhere” is a place you can start building toward now, without waiting for any number at all.

Which raises the obvious next question — the one the anxious professional is really asking underneath all of this. Maybe none of this applies yet. Maybe the disruption is overblown, the jobs aren't really going, and I'm worrying about a future that won't arrive. *Is this time even different?*

That's a fair challenge, and it deserves the most honest chapter in the book.

Sources for this chapter's empirical claims: Kela/VATT Finnish basic income results (2020); OpenResearch Unconditional Cash Study (2024); lottery-windfall wellbeing studies; the retirement-and-wellbeing literature. All logged and (where noted) verified in ../../evidence_base.md.

Chapter 3 — Are You Just Catastrophising?

Every previous time someone announced that machines were about to end work, they were wrong. The Luddites smashed the looms; textile employment, in the long run, grew. The mechanisation of farming threw most of humanity off the land; we did not end up with permanent mass unemployment, we ended up with cities and entirely new categories of jobs nobody in 1850 could have named. The ATM was supposed to end bank tellers; the number of tellers rose for decades afterward.

This track record is the optimist's strongest card, and it is a genuinely strong card. The honest version of the optimistic case is not "don't worry." It is: *every prior automation panic underestimated the economy's ability to invent new work, and the burden of proof is on anyone claiming this time is different.* That deserves to be met head on, not dismissed.

So: is this time different? The honest answer has to be split in two, because part of it we can reason about and part of it we simply have to watch.

A note on what kind of chapter this is. We're leaving the firm ground of data now — what follows is reasoning by analogy and mechanism, not measurement of a future that hasn't happened. My job here is not to tell you how it turns out. It's to hand you the cleanest way I know to *think* about it: to mark exactly where the reassuring historical analogy holds, and the specific places it might snap.

Where the analogy holds

At the level of structure, the AI transition really does rhyme with the Industrial Revolution, and we should give the optimists this much:

- Some kind of human labour gets automated.
- The old roles built around it shrink or vanish, often painfully, often fast for the people caught in it.
- There is a wrenching adjustment period.
- Eventually a new equilibrium forms — and crucially, *how humane that equilibrium is* depends enormously on institutions (unions, public schooling, the weekend, the welfare state) that, in the Industrial Revolution, arrived decades after the suffering started.

If the AI transition follows this template, the lesson is double-edged. Yes, there is probably a new equilibrium on the far side. But the Industrial Revolution's far side took the better part of a century to become livable, and the livability was *built*, not automatic. "It worked out in the end" is true and is cold comfort to everyone who lived and died in the middle.

The three places it might snap

Here is the part I want you to actually take away — not a verdict, but a **watchlist**. There are three specific variables on which the comforting analogy could break. The value of naming them is that each is, in principle, *observable*: instead of arguing about optimism versus pessimism in the abstract, we can watch these three dials.

1. What is being automated. This is the deepest one. Each prior wave automated one *layer* of human contribution and pushed people up to the next. Industrial machines took **muscle**; humans retreated to **routine cognition** — clerking, calculating, processing. Computers then took routine cognition; humans retreated again, to **judgement, creativity, and the handling of the non-routine**. Every time, there was higher ground to retreat to.

AI aims directly at that higher ground — at judgement, synthesis, creativity, the non-routine itself. The honest question is not "will AI take jobs" but: *if the highest cognitive ground is now contested, is there a next rung to retreat to, or is this the top of the ladder?* I don't know the answer. Nobody does. But notice that the historical reassurance secretly depended on there always being a higher rung — and that is precisely the assumption now in question.

2. Speed. The Industrial Revolution unfolded across *generations*. That timing did something easy to miss: the people whose livelihoods were destroyed were largely not the same people who staffed the new factories. The handloom weaver was ruined; his grandchildren grew up inside the new industrial world and found their place in it. The adjustment happened *between* generations, which is the only speed at which that kind of adjustment is bearable.

The mechanism that "rescued" us was real but slow. So the question is whether AI's diffusion moves at a generational pace or something far faster. If the capability and its deployment outrun the rate at which people can retrain, careers can turn over, and new institutions can form, then "in the long run it works out" never gets the long run it needs. The transition cost, normally smeared across decades, gets concentrated onto specific liv-

ing people.

3. Whether the “new jobs will appear” mechanism still works. This is the subtlest, and the most important. It is *true* that every previous wave created new categories of work. But look at *why* that rescued us: when the new jobs appeared, **humans were better at them than the machines of the day**. The new work was, by construction, work the current machines couldn’t do.

That precondition is exactly what AI calls into question. If a new category of valuable work appears and a general-purpose AI can also, immediately, do it — then the historical escape hatch, the thing that bailed us out every single previous time, quietly stops functioning. The optimist is not wrong that new work will appear. The open question is whether *we* will be the ones doing it.

What this chapter can and can’t tell you

Notice what I have and haven’t claimed. I have *not* told you AI will cause mass unemployment — that’s a forecast I’m not entitled to make. I’ve told you the three conditions under which the reassuring history would fail to repeat, and I’ve told you they’re observable.

So the honest stance for the anxious professional is neither panic nor the breezy “we’ve always been fine.” It’s: *watch these three dials*. Is the automation reaching genuine judgement and creativity, or stalling at the routine? Is it moving at generational speed or faster? And when new kinds of work appear, are humans doing them — or is the AI doing those too? You will, in your own field, often have a better early read on these than any pundit.

It would be a cheat, though, to leave you with three abstract dials and no early readings. As of 2026, the picture is genuinely mixed: AI’s diffusion is no longer in doubt (roughly 40% of US employees now use it at work), but on Claude.ai the usage still tilts toward *augmenting* human work over *automating* it (about 52% to 45%) — and the broad job market for exposed professions has, in aggregate, held up. There is one sharp exception, and it lands exactly where the analogy is most stressed. I owe you the full account of where the live data points and where it doesn’t — that is the next chapter.

Chapter 4 — The First to Hit the Wall

The last chapter ended on a frustration: to know whether “this time is different,” we have to look at what’s actually happening in the labour market right now — and the economists disagree.

So let me tell you what the live data does and doesn’t show, honestly, including the parts that cut against my own argument. Because the most interesting thing in it is not the headline. It’s *where* the damage is — and isn’t — already visible. And that turns out to be the hinge of this whole book.

The dog that isn’t barking

Start with the surprise, because intellectual honesty demands we lead with the evidence against alarm.

As of 2026, the broad white-collar catastrophe has **not** shown up in the aggregate data, and we should sit with how genuinely uncertain the picture is. AI adoption itself is no longer in doubt: Anthropic’s Economic Index, built from around a million real Claude conversations, finds roughly 40% of US employees now use AI at work, up from about 20% in 2023 — and that on Claude.ai as of late 2025, usage still tilts slightly toward *augmentation* over *automation* (about 52% to 45%). People are using it, heavily, and so far more to extend their work than to replace it. Meanwhile the aggregate job market for exposed professions has not cratered; by several measures employment and wages in AI-exposed industries have held up.

But there is one exception, and it is precise enough to take seriously. The Stanford Digital Economy Lab — Erik Brynjolfsson and colleagues, using high-frequency ADP payroll records covering millions of US workers — found that workers aged **22 to 25 in the most AI-exposed occupations** saw their employment fall by roughly **16% in relative terms** after generative AI spread, even as *older* workers in those very same jobs grew by 6–9%. The damage, where it shows up at all, is concentrated exactly where AI *automates* rather than augments — and it lands on the youngest, at the entry level. (Even this is contested: some economists note the entry-level dip began before ChatGPT and resist pinning it cleanly on AI. I’ll honour that uncertainty rather than wave it away.)

I have to give the optimist their due here. If you expected AI to be visibly gutting white-collar work by 2026, the aggregate data says: not yet, and maybe the adjustment really is being absorbed the way it was every time before. Anyone who tells you the white-collar apocalypse is *currently measurable* is getting ahead of the evidence. What we have is one narrow, real signal — the young, at the bottom rung — inside an aggregate that is otherwise holding.

But “not yet, in aggregate” is not the same as “nowhere.” There is one place where the future this book worries about is not a forecast at all. It has already, fully, happened. And it happened in plain sight.

Stack Overflow: the experiment that already ran

Come back to where the Prologue started. Stack Overflow — the global commons where programmers answered each other’s questions — lost roughly **45% of its monthly contribution in the first year after ChatGPT, and around 95% by early 2026**. Not a slowdown. A near-total collapse of a fifteen-year-old ecosystem of voluntary human contribution.

Why did this corner of the world fall so completely, while the broader white-collar job market didn’t? The answer is the most important idea in this chapter, and it’s structural:

Stack Overflow had no employment contract holding it together. Nobody was paid to answer questions. There was no manager, no salary, no notice period, no severance, no inertia of “I need this job to pay rent.” The only thing holding a contributor there was that the activity was *giving them something* — competence, status, community, the pleasure of being the one who knew. The moment a machine could supply the *asking* side better, there was nothing — no contract, no paycheck — buffering the *answering* side from simply evaporating.

A salaried job has enormous stickiness. Even when AI can do much of it, the contract, the organisational habit, the sheer friction of firing and reorganising all slow the unwinding to a crawl — which is exactly why the aggregate white-collar data still looks calm. Voluntary contribution has none of that stickiness. So when AI arrives, the post-work future shows up there **first, fast, and unbuffered**.

That makes Stack Overflow something rare and valuable: an **early laboratory of post-labour life**. It is a clean look at what happens when an activity that supplied people with meaning — but not income — becomes something a machine can do. The contracts that are currently disguising the effect everywhere else simply weren’t there to disguise it.

(I spent weeks inside that laboratory, studying exactly who stopped and who stayed. What I found surprised me and corrected several of my own confident guesses — but that’s a story for later in the book. For now, the headline is the collapse itself, and what it previews.)

Why the professionals are forerunners, not a special case

Now put the two pieces together — the calm aggregate, and the collapsed commons — and the shape of the argument appears.

The anxious professional reading this is not living in a different world from everyone else. They are living in everyone else’s *near future*, brought forward. They are closest to two edges at once: their work is among the most exposed to AI, **and** — especially for the young, the freelance, the gig-bound, the between-jobs — they often have the least contractual buffer protecting them from it. They are, in both senses, the first to hit the wall.

But I have to stop and steelman the strongest objection to this whole move, because it’s a good one and the honesty of the book depends on facing it. Maybe Stack Overflow didn’t collapse because it lacked a contract. Maybe it collapsed because it was a *perfectly substitutable special case*: asking a coding question and getting an answer is almost exactly what a chatbot does, end to end, with no residue. Very little white-collar work is substitutable that cleanly — most of it is tangled up with judgment, accountability, relationships, and context that a model can’t simply swallow whole. If that’s the real reason, then Stack Overflow isn’t a preview of everyone’s future at all; it’s just the one corner of knowledge-work that happened to map one-to-one onto what the machine does, and the rest of us are safer than its collapse makes us feel.

So I tried to settle it, rather than just worry about it in prose — and the attempt is worth telling you about, because it failed in an instructive way. The contract theory makes a sharp, checkable prediction. Take two people doing the same exposed work, equally substitutable by the machine — one salaried, one freelance. If the contract is what matters, the freelancer’s job should erode first and fastest, because nothing buffers it. So I built a small pre-registered study on three years of U.S. Current Population Survey data — froze the hypothesis before looking at the outcome, the discipline that keeps you honest — and compared, within the same occupations, the employment of unincorporated self-employed workers (no contract) against salaried employees (full contract), as a function of how exposed each occupation is to language models.

The honest result: I could not find the effect. Through early 2026, no-

contract workers in the most AI-exposed occupations were not losing jobs measurably faster than salaried workers in those very same occupations. The estimate leaned the way the theory predicts, but it was statistically indistinguishable from zero — and, tellingly, a placebo test run on the year *before* ChatGPT produced exactly the same faint lean, which means even that whisper of a signal wasn't ChatGPT's doing. It held no matter how I cut the data.

I want to be careful, because this cuts both ways. It does not refute the contract idea. The broad data are noisy, self-employment is measured crudely, and — most importantly — the AI employment shock is barely visible in this data *for anyone* yet, even salaried workers in exposed jobs, so there may simply be nothing for a contract to be buffering against so far. The test may be too early. But it does fail to confirm the idea exactly where I most expected it to show — and that costs me something. If the no-contract workforce were really being hit fast and unbuffered, three years of national data ought to be starting to show it, and they aren't. On this evidence the Stack Overflow collapse looks a little less like the leading edge of a broad no-contract unravelling, and a little more like what the objection warned: a uniquely substitutable special case. I have to let that pull my confidence down a notch, and I will.

So let me be exact about what kind of claim the next step is. The leap from “Stack Overflow collapsed” to “the professionals are forerunners for everyone” is not a measured finding — and it is now an inference with one failed confirmation behind it. It is *reasoning*: a bet that contracts are mostly *delaying* the same force, not permanently exempting most work from it. I still find it plausible, because the mechanism is sound and the early signals on the salaried side (the young, at the entry level) point that way — but after a direct test came back empty, I can no longer call it *more likely than not* with a straight face. It is a live, unsettled worry. Held that honestly, it is still worth making, because the stakes are high and the data are early. But you should hold it as a reasoned worry shadowed by a null result, not a demonstrated fact, and so will I.

With that fence built, here is the inference itself. This is why I've written a book that starts with programmers and translators and illustrators but is not, finally, about them. The contract-based stickiness that is currently keeping the aggregate numbers calm may well be a delay rather than a reprieve. What hit Stack Overflow without a buffer may yet reach the rest of work with one — slower, messier, cushioned by institutions — as the same underlying force, *if* the substitution objection above doesn't turn out to be the whole story. I no longer get to assume it won't. On that reading, the forerunners are seeing, early and undisguised, the question

that is eventually everyone's:

When the thing you produce can be produced without you — when your output no longer needs you — who are you, and where does your life get its meaning?

That is no longer a programmer's question, or a translator's. It is the question of the lottery winner and the retiree from Chapter 2, the question latent in Marienthal, the question every human in a wealthy automated society will eventually face in their own form. The professionals are just the first to be handed it in writing.

The forerunner's advantage

I want to end this chapter on the consolation, because it's real and it's earned.

Being first to hit the wall sounds purely like misfortune. It isn't. The forerunner has one thing the latecomer won't: **time, and warning**. You can see the question coming while you still have the income, the skills, the energy, and the years to do something about it. The retiree who declines is usually the one for whom the loss of work arrived all at once, unexamined, with no plan to replace what it had been quietly supplying. You are being given the opposite: the chance to see the loss clearly, name it (Chapter 1 gave you the list), understand why money won't cover it (Chapter 2), and start building the alternative *before* you need it.

The rest of this book is about that building. But first we have to be honest about the stakes — about how badly this can go at the level of a whole society when the wall is hit without preparation. That's the darkest chapter, and it's next.

Chapter 5 — How Badly Can This Go?

Every chapter so far has been tightening a worry. Work supplies five things money can't replace. Cash doesn't fill the gap. The reassuring history of automation might not repeat. And in the one place we've watched the future arrive without a buffer, the collapse was near-total. It would be dishonest to spend four chapters building that and then flinch from the obvious next question: *at the scale of a whole society, how bad can this get?*

This is the darkest chapter in the book, and also the one where it would be easiest to lie to you — in either direction. I could frighten you with a collapse the evidence doesn't support, or reassure you with a calm the evidence doesn't earn. I'm going to try to do neither, which means being unusually careful about the line between what we actually know and what we only fear.

The one thing we know is grim

Start with the firm ground, because it is solid and it is sobering.

For decades, the economists Anne Case and Angus Deaton tracked something they came to call **deaths of despair** — suicide, drug overdose, alcoholic liver disease. Among middle-aged American adults without a college degree, these deaths rose steeply as the stable work that had anchored their communities drained away over a generation. Between the late 1990s and the late 2010s, overdose deaths in this group rose more than four-fold; suicides and alcohol-related deaths climbed with them. Case and Deaton's argument is not that these people simply ran out of money. It is that they lost the whole structure that work had organised — the standing, the routine, the sense of a place in the world — and that the loss showed up, in the end, in the mortality statistics. When work and the meaning and community bundled with it collapse for a population, the result is not only sadness. It is earlier death.

It helps to picture the texture behind that statistic, because the number alone lets you keep it at arm's length. Think of one of the towns this actually happened in: built around a single mill or plant, its whole day once organised by the shift whistle — the same men clocking out together, filling the same bar at four in the afternoon, the union hall, the church softball league, the sense of being somebody who *did* something. Then the work

goes. A decade later the plant is a fenced lot, and the same men are home on a weekday afternoon with the television on and the hours unmarked, no reason to be anywhere by any particular time. It is the same scene Marienenthal described from the other side of an ocean and a century — people who slowly stopped knowing what day it was — only this time the slow unravelling shows up, eventually, in the overdose and suicide counts at the county coroner's office. That is what the curve is made of.

That is the strongest evidence we have that the stakes here are mortal, and I want it to land with its full weight. But I also have to fence it honestly, because this is exactly the kind of finding it's tempting to stretch. It is correlational, and badly confounded — the opioid epidemic, the American healthcare system, the specific wreckage of deindustrialisation in particular towns in a particular era all run through it. It tells us, unmistakably, that the loss of work-and- meaning *can* be lethal at population scale. It does not tell us that a future shaped by AI would follow the same path. That last step is a guess, and I'll keep treating it as one.

Where the frightening stories outrun the evidence

Past that single piece of hard ground, the darkness gets speculative fast — and here I owe you the discipline of *not* dressing up fear as knowledge.

There is, for instance, a famous and seductive parallel: the “behavioural sink.” In the 1960s the ethologist John Calhoun gave colonies of rats unlimited food and space but let their social roles collapse, and watched them descend into withdrawal, aggression, and the breakdown of normal behaviour and reproduction. As a metaphor for a society of abundance without purpose, it is almost irresistible — and it has been reached for, again and again, by people describing everything from falling birth rates to online life. I am deliberately not going to lean on it, and I want to be candid about why. It is a metaphor, drawn from rodents, and metaphors from rodents are not evidence about human societies. I know the pull of it firsthand: I once built an entire computational model on this very analogy, and eventually had to abandon it precisely because it could explain anything and therefore predicted nothing. I'd rather show you that scar than sell you the metaphor.

A second fear is more abstract but worth naming: if AI can out-produce everyone, what happens to *status* — which is, by nature, relative and zero-sum? It may simply migrate to new ground — authenticity, presence, being visibly human — but I have little hard evidence for how that plays out, and I'll flag it as the speculation it is.

And there is a third question, the one that actually decides whether “bad” becomes “catastrophic”: does meaning-loss *spread*? Does one person’s withdrawal quietly pull down the people around them, so that disengagement compounds through a community the way a contagion would? If it does, the downside is far worse than any individual tally suggests. Here, for once, I have a thread of real evidence — and, interestingly, it seems to cut *against* the scary version. In my own study of Stack Overflow, the community’s underlying motivational structure stayed remarkably *stable* straight through the shock; the people who left didn’t appear to detonate a cascade in those who remained.

But I have to be harder on that reassurance than my instincts want to be, because the honesty of this whole book lives exactly here — in not over-trusting evidence just because it points the way I’d like. The truth is I don’t fully trust even this reverse-evidence, and I’ll tell you why. My measure of that community’s “structure” was built from *behaviour* — who comments, who edits, who answers — and behaviour like that mostly captures people’s stable dispositions, not their inner reaction to a shock. So a finding of “structure stayed stable” might just mean my instrument was never able to see the reaction in the first place. I’m not guessing at that weakness after the fact: I built a placebo test into the study precisely to catch myself, and it caught exactly this — a pattern I’d been ready to read as a real effect turned out to be the tool measuring temperament, not change. The stability I found is therefore weak comfort. It is genuinely some evidence against the contagion story, and honesty runs in both directions: when the one piece of evidence I have leans away from the frightening version, I’m obliged to say so. But I hold it loosely. The instrument that produced it may simply have been blind to the very thing we most fear.

The honest shape of the danger

So here is what I can actually claim, with neither flinch nor exaggeration.

The risk is real, and the stakes — the deaths of despair tell us this much — can be mortal. But the *shape* the danger takes is not fixed in advance. The worst outcomes in history were never automatic; they happened in the specific places where the adjustment came fast, hit people without a buffer, and was met with no institutional response at all. That is the crucial distinction, and it’s the hinge of everything that follows in this book: this is a **warning, not a prophecy**. A prophecy you can only await. A warning you can act on.

And that distinction is also where I can offer you something, personally, even in the bleakest chapter. The honest takeaway here is *not* “society will

collapse, so despair.” It is the opposite. The downside is serious enough to take seriously — and it is conditional, hanging on factors that can still be changed, which is the entire reason that preparing, both in your own life and collectively, is worth the effort. Despair is the one response the evidence specifically does *not* license, because the bad endings are contingent, not fated.

The rest of the book is about that contingency: where a person can still find meaning when work recedes, and how a society might steward AI so that more people get the chance to. We turn to the first of those next.

Chapter 6 — Finding Your Place Again

This is the chapter the whole book has been walking toward. Everything before it was, in a sense, clearing the ground: showing what work was quietly giving you (Ch.1), why money won't replace it (Ch.2), why the comforting history might not repeat (Ch.3), and why the people hit first are forerunners for everyone (Ch.4–5). All of that is diagnosis. This chapter is about what you actually *do* — and it's the part I most want to get right, because if this book helps you at all, it will be here.

I won't pretend to have your answer. Finding your place again is something only you can do, in your own life, with your own materials. But I've been doing it myself, and I can offer something honest: not a map of where you should end up, but a clear look at one real path through — mine — and the turn at the center of it that I think matters for almost anyone.

Let me just tell you what happened.

A small, true story

I wanted to make a game. I'm not a game developer — it was never my field. A few years ago, “make a game” would have meant years of learning before I could even start, so I never started.

Then AI got good enough that I could. The model could write code I couldn't, fill gaps I didn't have the training to fill, take me from *can't* to *can* in a fraction of the time. That is the first thing AI can be in your life, and it's real: **a lever**. Each time the models improve, some project that was impossible for you quietly becomes possible. (This very book, and the research behind it, is another example — an independent person doing analysis that would have needed a team a few years ago.) Positioning yourself so that every leap in AI *expands* what you can do, rather than threatening what you already do, is a genuine reorientation. It's worth taking seriously.

But here's the part I have to be honest about, because it's where the easy “just use AI!” advice falls apart.

The same lever that let *me* make a game lets *everyone* make a game. The thing AI handed me, it handed to millions. The moat — the years of spe-

cialized skill that used to make game-making scarce and therefore valuable — AI just drained that moat, for me and for everyone at once. So if I'd been doing this to build something commercially valuable, to be one of the few who *could*, I would have arrived to find I was one of millions. The market value of “I can make a game” is collapsing precisely *because* the lever works for everyone.

This is the trap underneath all the cheerful “embrace AI” advice, and most of it never says so. AI as a lever and AI as a moat-destroyer are the *same thing*. You don't get one without the other.

The turn

So here is the turn — the one real thing I want to hand you, and it's not specific to games or to me.

When I noticed the moat was gone, I had a choice about where my sense of *why I'm doing this* was anchored. I could keep it anchored to the **outcome** — the commercial value, the status of being one of the few who could, the hope of an audience or a payoff. If I did that, AI was going to make me more and more anxious, because every one of those things was getting scarcer and less in my control. Or I could let it rest somewhere else: on the thing I'd actually been enjoying without quite admitting it — **the process itself**. Going from *can't* to *can*. The pleasure of making the thing, watching it come alive under my hands, learning a domain I'd been shut out of.

That second place — the process, the learning, the doing — is something AI cannot take from you, and the reason it can't is important. It's *not* because it's some hidden fortress AI hasn't reached yet. It's because it was never out in the market in the first place. Meaning anchored to an external result — money, status, recognition, being ahead of others — was always hostage to things you don't control, and AI is now yanking hard on exactly those things. Meaning anchored to what you can control — your own growth, the experience of doing, the people right in front of you — was never for sale, so there's nothing for AI to devalue.

It's worth giving this choice a name, because once you have it you'll start seeing it everywhere. Call it **outcome-anchoring versus process-anchoring**: whether the thing that makes your effort feel worth it sits in the *result* you're chasing — the money, the status, the win — or in the *doing* itself — the learning, the making, the people you do it with. It isn't a slogan and it isn't a personality type; it's just a question about where you've set the anchor, and you can move it. Almost everything practical in this book comes down to that one move.

This isn't resignation. I'm not telling you to give up on outcomes and be content with crumbs. I'm telling you something more practical: **in a world where AI is systematically driving down the scarcity value of almost any output, anchoring your meaning to outcomes is a bet that gets worse every year — and anchoring it to process is a bet that gets safer.** It's not a consolation prize. It's the more clear-eyed allocation of where to place the thing that makes your days feel worth living.

And here I can offer more than my own experience, because psychologists have actually measured this. Decades of research within Self-Determination Theory (Kasser, Ryan, Deci and colleagues, replicated across cultures and age groups) find that people who orient their lives around *extrinsic* goals — money, status, image, the approval of others — report **lower** wellbeing and more distress than people oriented toward *intrinsic* goals: growth, relationships, contribution. That much you might have guessed. But one study cuts deeper and is worth pausing on. Niemiec, Ryan and Deci followed people through the year after college and looked not at what they *pursued* but at what they *attained*. The people who achieved their intrinsic goals got the wellbeing boost you'd expect. The people who achieved their **extrinsic** goals — who actually got the money, the status, the recognition they were chasing — got **no boost at all**, and if anything reported *more* ill-being. Read it again: *succeeding* at the external goals didn't deliver. The prize, even when won, was empty of the thing they wanted from it. I'll be precise about its weight: this is a single longitudinal study, not a settled law, so I lean on it as a sharp illustration rather than proof — but its direction lines up with a large body of Self-Determination research pointing the same way, which is why I trust the arrow even if I won't oversell any one study.

That finding mattered before AI. What AI does is press on exactly the same nerve from the outside: it is driving down the value of the external prizes *and* the evidence says those prizes wouldn't have made you whole even if you'd kept them. Both blades of the scissors point the same way — toward moving your weight onto what was load-bearing all along.

The thing people will get wrong

There's a tempting, and dangerous, misreading of all this, and I want to head it off because I nearly fell into it myself.

You might think: *fine — I'll relentlessly use AI to go from can't to can across every domain, get fast at learning anything, and then when a real opportunity appears I'll be faster than everyone and cash in.* And it's true that the ability to go from *can't* to *can* quickly is the one skill that transfers across every field

— far more durable than any specific expertise AI can absorb. Breadth, and speed of learning, are genuinely valuable.

But notice what just happened in that sentence: I smuggled the outcome back in. “Cash in when the opportunity comes” puts your meaning right back on an external result — one that’s *possible but not in your control*. Whether the opportunity comes, whether you’re faster than the next person — that’s not yours to decide, and making it your goal just rebuilds the anxiety you were trying to escape, one level up.

So hold it the other way around. Range widely, get fast at going from *can’t* to *can* — but because the going-from-can’t-to-can is *itself* the good part, not because it’s an investment in a future payoff. If the payoff comes, it’s a bonus. If it doesn’t, you’ve lost nothing, because the process already paid you. That’s the whole difference between a life AI makes more anxious and one it makes richer, and it turns entirely on where you let your meaning rest.

Other places to stand

The game story is one path — the one I know from the inside. It’s not the only shape “finding your place again” can take, and I don’t want to pretend my route is yours. A few others, more briefly, each a variation on the same move of anchoring meaning where AI can’t reach because it was never in the market to begin with:

- **From producing to discerning.** When AI can generate infinite output, what stays scarce — and satisfying — is *judgment*: knowing what’s worth making, what’s good, what matters. For some people the reorientation isn’t doing more, it’s developing taste and direction.
- **Toward the people in front of you.** The most durable anchor of all, and the one with the strongest evidence behind it: care, presence, being someone specific people actually rely on. AI doesn’t compete here, because the value was always in a *human* being the one there — and this is now more than a hunch. Across nine experiments with over 6,000 people, the *identical* caring words were rated more empathic and more satisfying when believed to come from a human than from an AI; merely disclosing AI authorship lowered the felt empathy. The thing care gives is partly constituted by a human being its source — which is exactly the part AI can’t supply.
- **Redefining “enough.”** Not every adaptation is finding a new place to *strive*. For some it’s moving the weight of meaning off work and output altogether — toward family, community, craft, faith, the

plain experience of a day. This is the oldest wisdom there is, and AI's pressure may simply be making it legible again.

I'll be plain about what these are, and what they aren't. The structural claim — that meaning rooted in *what you control and experience* is safer from AI than meaning rooted in *outcomes and output* — is something I think the evidence and the reasoning in this book genuinely support. But which of these paths is right *for you*, or whether it's one I haven't thought of, is not something I can hand you, and I won't pretend otherwise. That part is yours. What I can do — what this whole book is trying to do — is help you see the choice clearly, so that you're making it on purpose instead of having it made for you.

What this asks, and what it can't reach

I have to add one honest limit, because this path is not equally open to everyone, and pretending otherwise would be exactly the kind of dishonesty this book refuses.

Everything in this chapter assumes you have *something* to reorient toward — curiosity, some inner pull toward a project, the energy and security to go from can't to can. That's not nothing to assume. For someone whose work was their livelihood, lost suddenly mid-life, with bills due, "enjoy the process of learning a new domain" can be a luxury they can't afford this month. The individual reorientation this chapter describes works best for people who already have a measure of agency and slack — and saying so honestly is important, because it's exactly where the individual path runs out and something larger has to take over. That's the next chapter: not what *you* do, but what a society owes the people for whom "just adapt" isn't enough on its own.

Chapter 7 — What Policy Can Do

The last chapter ended where the individual path runs out. “Find your place again” works best for people who already have some agency and slack — curiosity, energy, a project, a buffer. For the person whose livelihood vanished suddenly with bills due, “enjoy the process of learning” is not yet advice; it’s a luxury. Something larger has to catch them. That something is policy, and this chapter is about what it can honestly be asked to do.

I’ll be upfront about two things first, because they set the whole register.

First, this is the one chapter where you, the individual reader, are mostly not the agent. Everything else in this book has been about what you can see, understand, and do. Policy is not that. It happens at the level of states and institutions, on timescales and through mechanisms that no individual controls. So I’m not going to pretend this chapter is a to-do list for you. It is here for a different reason: so you can understand the environment you’re adapting *into* — read it the way a sailor reads weather, not the way a captain gives orders.

Second, the frame is not “fight AI.” I want to be clear about this because the obvious instinct — and a lot of policy talk — treats AI as a threat to be walled off, contained, slowed. That’s the wrong frame, and not just tactically. AI is going to arrive and stay, and it may bring enormous good. The better question is not *how do we hold it back* but *how is its benefit and its harm distributed among ordinary people?* So the right center for policy, the thing everything else should orbit, is simple to state:

Amplify the benefit AI brings to citizens; reduce the likelihood and severity of the harm it does to them.

Not resistance. Stewardship. Everything below is a variation on that one sentence — and none of it is a prediction. These are design and value choices, and the honest test for each is not “is it true” but “is it a robust bet — worth doing across a wide range of how AI actually unfolds?”

Reducing the harm

The harms are the more urgent half, because they fall hardest and fastest on exactly the people with the least slack — the ones Chapter 6’s individ-

ual path can't yet reach.

- **A floor, honestly understood.** Chapter 2 settled what cash can and can't do: it buys relief from material stress, not meaning. That makes income support — UBI or otherwise — necessary but not sufficient. As policy, the honest version is: *a floor that prevents destitution, while nobody pretends it has solved the deeper problem.* Funding the floor and declaring victory would be funding Marienthal without curing it.
- **Cushioning the speed.** Recall from Chapter 3 that what made past transitions bearable was *time* — the damage spread across generations. The cruelest version of an AI transition is a fast, unbuffered one concentrated on specific living people. A great deal of what policy can helpfully do is simply *slow the landing for those caught underneath it*: transition support, time to reorient — and the obvious answer here is more complicated than it looks. The reflexive policy response to displacement is “retrain them.” But the evidence on retraining is genuinely sobering: large US evaluations of government training programs (WIA/WIOA, Trade Adjustment Assistance) have repeatedly found little or no improvement in earnings or employment versus a control group — in some cases participants did *worse*. That is the measured record, not my skepticism talking. There is one honest exception worth naming, because it points at what does work: training tied to a *specific employer or a concrete job waiting at the end* — sectoral and employer-linked programs — has a markedly better record than the generic “go learn new skills and hope” model. So “just retrain people” is not the solution it's usually waved around as — though “retrain people into a job that actually exists, with the employer in the room” is a real exception. That doesn't mean abandon transition support; it means be honest that cushioning the fall is mostly about *time, income, and dignity* during the adjustment, not a confident promise that we can re-skill everyone into the next good job. Not slowing AI — slowing the fall of the people it displaces, without overpromising about where they land.
- **Distributing the gains.** AI's productivity dividend is real and growing, and the fiscal picture here is worth being concrete about. A meaningful UBI is expensive — Andrew Yang's \$1,000-a-month proposal was costed at roughly \$2.8–3 trillion a year, which the current US tax structure simply can't carry without major reform. The standard rejoinder is “fund it from the AI surplus” — via sovereign wealth funds, public equity stakes in AI firms, or land-value taxes. The clearest working precedent is small but real: since 1982 every Alaskan, including children, has received an annual check from oil royalties pooled into a sovereign fund — usually \$1,000–\$2,000, oc-

casionally more — credited with measurably reducing poverty there without measurably cutting work. It is the proof of concept that a citizen dividend funded from a natural resource *can* run for decades; by one estimate AI need only reach about 5–6× today’s automation productivity to finance a UBI worth ~11% of GDP on the same logic. That may well become true. But notice where the surplus goes *first*: in late 2025, the big AI firms’ capital spending was running around \$142 billion a *quarter* — vastly more new money flowing to capital than any legislature has moved toward citizens. The dividend accrues to whoever owns the models first, and to wages later and unevenly. So “can we afford to share AI’s gains” is not really a technical question. It’s a political one — about who captures the surplus — and it’s where “reduce the harm” and “amplify the benefit” meet.

Amplifying the benefit

The benefit side is the half that usually gets left to markets, but policy shapes it too.

- **Lowering the barrier to the lever.** Chapter 6’s lever — AI letting an ordinary person go from *can’t* to *can* — is currently most available to people who already have the access, the literacy, and the confidence to use it. The most democratizing thing policy could plausibly do is make that lever genuinely available to everyone: access, education, the basic capability to turn AI into agency. I should mark this one’s status carefully: that broader access *would* expand agency for those currently shut out is reasoning, not a measured fact — the rigorous evaluations of past technology-access programs are still coming in. But it is the bridge back to Chapter 6, because it widens the door to the individual path for the people currently shut out of it.
- **Protecting the spaces where humans matter to each other.** The most durable sources of meaning (Ch.6: care, presence, the people in front of you) are also the ones markets chronically underfund. Policy that resources care work, community, and the institutions where people are genuinely needed isn’t nostalgia — it’s investing in exactly the meaning-sources that survive AI, on behalf of the people who can’t privately afford to build them.

Why “robust” is the right standard

I can’t tell you AI will definitely cause mass displacement, so I can’t justify any of this with “it’s coming, prepare.” The defensible justification is the robust-bet logic: these are choices that look good across many futures.

And this isn't a rhetorical dodge — it's an actual, named method. Analysts at RAND (Robert Lempert and colleagues) formalized it as **Robust Decision Making**: when the future is *deeply* uncertain, you stop trying to predict the one true outcome and instead stress-test each option across a wide range of plausible futures, keeping the ones that hold up broadly. "Make good decisions without predictions" is a real discipline, used for problems like climate and water policy where, as here, no one can forecast and stakeholders don't even agree on the assumptions. That is exactly the standard I'm applying to AI policy.

By that standard: if AI displaces enormous amounts of work, a floor, a cushion, distributed gains, and a widely available lever are essential. If the disruption stays mild, they're still good — just less urgent. There's no plausible future in which "ordinary people share more of AI's benefit and absorb less of its harm" turns out to be the wrong thing to have aimed at. That asymmetry is what makes it a stewardship worth committing to rather than a gamble on a forecast.

Back to you

So read this chapter as context, not instruction. You don't set policy. But knowing what policy *could* do — and what it would take to push it toward stewardship rather than either denial or panic — is part of seeing your situation clearly. And there's a quiet handoff between this chapter and the last one: the better the floor, the cushion, and the access to the lever, the more people get to reach the individual path of Chapter 6 — the move from outcome-anchoring to process-anchoring — at all. Policy's deepest job, in the end, is to widen the doorway through which people can go and find their place again.

That's nearly the whole argument. What remains is to gather it up — and to mark, one last time, what we know, what we're guessing, and how to live in the space between.

Epilogue — Living Honestly With Uncertainty

I can't tell you how this ends. That isn't modesty; it's the truth — and a book that pretended otherwise would have broken, on its last page, the one promise it made on its first.

But across these chapters I've tried to show that “we can't predict it” is very far from “we can say nothing” — and even further from “do nothing.” Let me gather what we *can* stand on.

What we actually know — the firm ground

- Work was quietly supplying five things beyond income — structure, contact, purpose, status, competence — and losing them, not just the paycheck, is what does the real damage. (Ch.1)
- Money buys relief from material stress, not meaning. Every cash experiment we have points the same way. (Ch.2)
- In the one place AI hit voluntary human contribution without the buffer of an employment contract — Stack Overflow — the collapse was near-total. (Ch.4)

What we can reasonably infer — the careful reasoning

- The reassuring history of automation might not repeat, and there are three observable dials that will tell us — what's being automated, how fast, and whether humans do the new work. (Ch.3)
- The professionals hit first *may well be* forerunners rather than a special case; contractual stickiness *may well be* delaying, not preventing, the broader arrival of the same question — though a pre-registered direct test of that mechanism came back null, so this is a reasoned worry shadowed by that null, not a finding. (Ch.4)
- Meaning sourced in *being-human-doing-it* — care, craft, presence — is structurally more defensible against AI than meaning sourced in output. (Ch.6)

What we have to choose — the part that is ours, not the data's

- At the level of society: whether to steward AI so it amplifies what citizens gain and softens what they lose — rather than fighting it, or leaving the gains to whoever owns the models. (Ch.7)
- And at the level of a single life: where to anchor meaning when output no longer defines us. (Ch.6)

Two things to carry out

I want you to leave with two things, and I want to be honest about what each one is.

The first is a **tool**, and it's the solid kind — the habit of asking, of any claim about the AI future, your own included: *which kind of claim is this?* Is it something we genuinely know, something we can reasonably infer, or something someone is choosing and dressing up as fact? That single question is the best defence I know against the confident nonsense this subject attracts. It will not tell you the future. It will keep you from being fooled about it, which is more useful.

The second is a **stance**, and I'll be honest that this one is my counsel, not my finding — I'll keep that line bright even here at the end. It is the one move this whole book turns on, the one I named in Chapter 6: shift from outcome-anchoring toward process-anchoring. Begin, now, to loosen your sense of meaning from your output and your job title — from the result you're chasing — and re-anchor it in the doing itself and in the things this book's evidence and reasoning keep pointing back to: being needed, being present, connection, craft, the people and the place that are specifically yours. I can't prove that's the answer. I can tell you it's where the honest version of the argument leads, that it's a robust bet across the futures we can't choose between, and that it's the bet I'm making with my own life.

The forerunner, again

We started with the anxious professional watching AI move into their craft, and we end with everyone, because the professional was only ever the first to be handed a question that's coming for all of us: *when what you make can be made without you, who are you?*

Here is the consolation I actually believe, stated at its true weight. The premise hidden inside that frightening question — that you *are* what you produce — was never true to begin with. It was a convenient fiction of an

economy that needed your output, and so taught you to look for yourself in it. AI is calling that fiction's bluff: brutally, and perhaps in the end mercifully. What is left when output no longer defines you is not nothing. It may turn out to be more *you* than the old bargain ever let you be.

That is not a prediction, and I won't dress it as one. The future stays unknown. So I end where I am allowed to end: not with an answer, but with a way of standing in front of the question without flinching and without lying — clear about the little we know, honest about the much we don't, and unwilling, in either direction, to mistake one for the other. If this book leaves you even slightly better equipped to do that — and a little less alone in it — it has done what it set out to do.

This book grows out of an open, ongoing project. Its evidence, its data, the studies that proved its author wrong, and the running list of what it still needs to find out, are all public. If you want to hold me to the promise I made on page one — that's where to do it.

References

- Anthropic. (2026). *The anthropic economic index: Fifth report* (tech. rep.). Anthropic.
- Brynjolfsson, E., Chandar, B., & Chen, R. (2025, November). *Canaries in the coal mine? six facts about the recent employment effects of artificial intelligence* (Working Paper). Stanford Digital Economy Lab.
- Brynjolfsson, E., Li, D., & Raymond, L. R. (2023). *Generative AI at work. NBER Working Paper*, (31161).
- Case, A., & Deaton, A. (2020). *Deaths of despair and the future of capitalism*. Princeton University Press.
- Cesarini, D., Lindqvist, E., Östling, R., & Wallace, B. (2016). Wealth, health, and child development: Evidence from administrative data on Swedish lottery players. *Quarterly Journal of Economics*, 131(2), 687–738.
- Jahoda, M. (1982). *Employment and unemployment: A social-psychological analysis*. Cambridge University Press.
- Jahoda, M., Lazarsfeld, P. F., & Zeisel, H. (1933). *Marienthal: The sociography of an unemployed community* [English translation published 1971, Aldine-Atherton].
- Kangas, O., Jauhiainen, S., Simanainen, M., & Ylikännö, M. (2020). *Suomen perustulokeilun arviointi (evaluation of the finnish basic income experiment)* (Report). Ministry of Social Affairs and Health, Finland.
- Kasser, T., & Ryan, R. M. (1993). A dark side of the American dream: Correlates of financial success as a central life aspiration. *Journal of Personality and Social Psychology*, 65(2), 410–422.
- Lempert, R. J., Groves, D. G., Popper, S. W., & Bankes, S. C. (2006). A general, analytic method for generating robust strategies and narrative scenarios. *Management Science*, 52(4), 514–528.
- Lempert, R. J., Popper, S. W., & Bankes, S. C. (2003). *Shaping the next one hundred years: New methods for quantitative, long-term policy analysis*. RAND Corporation.
- Meaning Project. (2026a). *No contract-buffer effect on AI-exposed employment: A pre-registered triple-difference on CPS 2021–2026* [Project repository: `contract_buffer_study/PAPER.md`].
- Meaning Project. (2026b). *The Stack Overflow motivational-structure null: A pre-registered test under the ChatGPT shock* [Project repository: `so_sdt_study/PAPER.md`].

- Niemiec, C. P., Ryan, R. M., & Deci, E. L. (2009). The path taken: Consequences of attaining intrinsic and extrinsic aspirations in post-college life. *Journal of Research in Personality*, 43(3), 291–306. <https://doi.org/10.1016/j.jrp.2008.09.001>
- Noy, S., & Zhang, W. (2023). Experimental evidence on the productivity effects of generative artificial intelligence. *Science*, 381(6654), 187–192.
- OpenResearch. (2024). *Unconditional cash study: Three-year findings on \$1,000/month* (tech. rep.). OpenResearch Lab. <https://openresearchlab.org/>
- Paul, K. I., & Moser, K. (2009). Unemployment impairs mental health: Meta-analyses. *Journal of Vocational Behavior*, 74(3), 264–282. <https://doi.org/10.1016/j.jvb.2009.01.001>
- Rubin, M., Li, J., Zimmerman, A., Ong, D. C., Goldenberg, A., & Gross, J. J. (2025). Comparing the value of perceived human versus ai-generated empathy [Nine experiments, N > 6,000; AI disclosure reduces felt empathy]. *Nature Human Behaviour*.
- Tax Foundation. (2019). *Analysis of andrew yang's "freedom dividend" proposal* (tech. rep.). Tax Foundation. <https://taxfoundation.org/>